

Pandas

Data Science 2 / Data & AI 3

Revision

Revision - Indexing

- What is the proper indexing to retrieve the the values in the yellow squares?
- The blue squares?
- The red squares?

1	2	3	4	5
6	7	8	9	10
11	12	13	14	15
16	17	18	19	20
21	22	23	24	25
26	27	28	29	30

Revision

- What is the type of this numpy array?

```
X= np.array(  
    [[1,2,3],  
     [4,"5",6],  
     [7,8,9]  
    ])
```

Revision

- How to replace items that satisfy a condition without affecting the original array?

The input is: `array([0, 1, 2, 3, 4, 5, 6, 7, 8, 9])`

The expected output is: `array([0, -1, 2, -1, 4, -1, 6, -1, 8, -1])`

Agenda

1. Introduction to Pandas
2. Indexing and Selection
3. Operations and Missing Values
4. Merge and Join
5. Aggregation and Grouping
6. Working with Strings
7. Working with Time Series
8. Reading files





Introduction to Pandas

What is Pandas

Python library with flexible data structures developed for Data Scientists

DataFrame

Series

Data Structures are build on Numpy arrays

Series

Series

DataFrame

apples		oranges		apples		oranges	
0	3	0	0	0	3	0	0
1	2	1	3	1	2	3	3
2	0	2	7	2	0	7	7
3	1	3	2	3	1	2	2

What is Pandas

Importing exporting and processing multiple data sources




Uniform handling missing data

N.A.

Explicitly defined indexes enabling advanced indexing, slicing and subsetting

Time series functionality



Advanced data manipulation

- GroupBy
- Joining
- ...

Pandas Series

Series as generalized NumPy array

- Numpy array: implicitly defined integer index
- Pandas Series: explicitly defined index

```
data = pd.Series([0.25, 0.5, 0.75, 1.0], index=[2, 5, 3, 7])
```

```
data[5] # 0.5
```

Series as specialized dictionary

- Python dictionary: values can have different types
- Pandas Series: all values have the same type (efficiency!)

```
population = pd.Series({'be': 10, 'nl': 8})
```

```
data['be'] # 10
```

Pandas Dataframes

Dataframe as generalized 2D NumPy array

```
countries = pd.DataFrame([{'population': 11.7, 'area': 30688},  
                           {'population': 17.7, 'area': 41850}])  
countries
```

	population	area
0	11.7	30688
1	17.7	41850

Series as specialized dictionary

```
population = pd.Series({'be': 11.7, 'nl': 17.7})  
area = pd.Series({'be': 30688, 'nl': 41850})  
countries = pd.DataFrame({'population': population, 'area': area})  
countries['area']
```

or countries.area

	population	area
be	11.7	30688
nl	17.7	41850

Notebook and Exercise time!

Notebook


See 03.01-Introducing-Pandas-Objects.ipynb

Exercise time!

See 03.01_EX.ipynb



Indexing and Selection



Indexing and Selection

1. Data Selection in Series

- Series as dictionary
- Series as one-dimensional array
- Indexers: loc, iloc, and ix
 - Avoid *ix* because it is no longer available in modern pandas versions

2. Data Selection in DataFrame

- DataFrame as a dictionary
- DataFrame as two-dimensional array
- Additional indexing conventions

Indexing and Selection - Series

```
data = pd.Series(['a', 'b', 'c'], index=[1, 3, 5])
```

1	a
3	b
5	c

```
# explicit index: .loc
```

```
data.loc[1] # 'a'
```

```
# slicing
```

```
data.loc[1:3] # 1 a
```

```
# 3 b, explicit index: final index is included
```

```
# implicit index: .iloc
```

```
data.iloc[1] # 'b'
```

```
# slicing
```

```
data.iloc[1:3] # 3 b
```

```
# 5 c, implicit index: final index is excluded
```

```
# masking and fancy indexing
```

```
data[(data == 'a') | (data == 'b')]
```

```
data.loc[[1,3]]
```

Indexing and Selection - Dataframe

```
data= pd.DataFrame([ {'population': 11.7, 'area': 30688},  
                     {'population': 17.7, 'area': 41850}], index=['be', 'nl'])
```

```
data['density'] = data['pop'] / data['area']
```

	population	area	density
be	11.7	30688	0.000381
nl	17.7	41850	0.000423

implicit index: .iloc

```
data.iloc[:1, :1]
```

implicit index: final index is excluded -> 1x1

explicit index: .loc

```
data.loc[: 'nl ', : 'area']
```

explicit index: final index is included -> 2x2

with masking and fancy indexing

```
data.loc[data.population>15, ['area', 'density']]
```


Notebook and Exercise time!

Notebook


See 03.02-Data-Indexing-and-Selection

Exercise time!

See 03.02_EX.ipynb



Operations and Missing Values in Pandas



Operating on Data in Pandas

Operations in Pandas

Ufuncs: Index Preservation

Ufuncs: Index Alignment

Operations Between DataFrame and Series

```
df = pd.DataFrame(rng.randint(0, 10, (3, 4)),  
                  columns=['A', 'B', 'C', 'D'])  
df
```

	A	B	C	D
0	6	9	2	6
1	7	4	3	7
2	7	2	5	4

```
np.sin(df * np.pi / 4)
```

	A	B	C	D
0	-1.000000	7.071068e-01	1.000000	-1.000000e+00
1	-0.707107	1.224647e-16	0.707107	-7.071068e-01
2	-0.707107	1.000000e+00	-0.707107	1.224647e-16

Missing values

1. Handling Missing Data

2. Trade-Offs in Missing Data Conventions*

3. Missing Data in Pandas*

- `None`: Pythonic missing data*
- `NaN`: Missing numerical data*
- NaN and None in Pandas*

4. Operating on Null Values

- Detecting null values
- Dropping null values
- Filling null values

* Reading for context suffices

Missing Values

Pandas treats None and NaN as essentially interchangeable for indicating missing or null values

```
df = pd.DataFrame([[1, np.nan, 2],  
                  [2, 3, 5]])
```

detecting null values

```
df.isnull()
```

```
df.notnull()
```

dropping null values

```
df.dropna()
```

```
df.dropna(axis='columns', thresh=3)
```

filling null values

```
df.fillna(0)
```

drops rows

drops columns, with min 3 NAs

fill with NAs with 0

	0	1	2
0	1.0	NaN	2
1	2.0	3.0	5

	0	1	2
0	False	True	False
1	False	False	False

	0	1	2
1	2.0	3.0	5

Notebook and Exercise time!

Notebook

See 03.04-Missing-Values.ipynb

See 03.04-Missing-Values.ipynb


Exercise time!

See 03.03_EX.ipynb

See 03.04-EX.ipynb



Merge and Join



Merge and Join

1. Combining Datasets: Merge and Join
2. Relational Algebra
3. Categories of Joins
 - One-to-one joins
 - Many-to-one joins
 - Many-to-many joins
4. Specification of the Merge Key
 - The ``on`` keyword
 - The ``left_on`` and ``right_on`` keywords
 - The ``left_index`` and ``right_index`` keywords
5. Specifying Set Arithmetic for Joins
6. Overlapping Column Names: The ``suffixes`` Keyword
7. Example: US States Data

Merge and Join

```
df1 = pd.DataFrame({'employee': ['Bob', 'Jake'],  
                    'group': ['Acc', 'Eng',]})  
df2 = pd.DataFrame({'employee': [ 'Jake', 'Bob'],  
                    'hire_date': [ 2012, 2008]})
```

df1			df2		
	employee	group		employee	hire_date
0	Bob	Acc	0	Jake	2012
1	Jake	Eng	1	Bob	2008

```
# merge detects common column
```

```
pd.merge(df1, df2)
```

```
# can merge one-to-one, one-to-many, many-to-many
```

	employee	group	hire_date
0	Bob	Acc	2008
1	Jake	Eng	2012

```
# merge with different column names
```

```
pd.merge(df1, df3, left_on="employee", right_on="name")
```

```
# merge on index
```

```
df1a.join(df2a)          # same as pd.merge(df1a, df2a, left_index=True, right_index=True)
```

```
# merge on index and column
```

```
pd.merge(df1a, df3, left_index=True, right_on='name')
```

Merge and Join

default is 'inner' join

'outer', 'left', and 'right' joins

```
pd.merge(df6, df7, how='outer')
```

overlapping column names

```
pd.merge(df8, df9, on="name", suffixes=["_L", "_R"])
```

Notebook and Exercise time!

Notebook

See 03.07-Merge-and-Join.ipynb

Exercise time!

See 03.07-EX.ipynb



Aggregation and Grouping

Aggregation and Grouping

1. Aggregation and Grouping
2. Planets Data
3. Simple Aggregation in Pandas
4. GroupBy:
 - Split, apply, combine
 - The GroupBy object
 - Column indexing
 - Iteration over groups
 - Dispatch methods

Aggregation and Grouping

4. GroupBy: Split, Apply, Combine

- Aggregate, filter, transform, apply
 - Aggregation
 - Filtering
 - Transformation
 - The apply() method
- Specifying the split key
 - A list, array, series, or index providing the grouping keys
 - A dictionary or series mapping index to group
 - Any Python function
 - A list of valid keys
- Grouping example

Simple Aggregation

```
df = pd.DataFrame({'A': [1, 2, 3],  
                  'B': [3, 4, 5]})
```

	A	B
0	1	3
1	2	4
2	3	5

```
df.mean()
```

A	2.0
B	4.0

```
df.mean(axis='columns')
```

0	2.0
1	3.0
2	4.0

```
df.describe()
```

	A	B
count	3.0	3.0
mean	2.0	4.0
std	1.0	1.0
min	1.0	3.0
25%	1.5	3.5
50%	2.0	4.0
75%	2.5	4.5
max	3.0	5.0

Aggregation	Description
<code>count()</code>	Total number of items
<code>first()</code> , <code>last()</code>	First and last item
<code>mean()</code> , <code>median()</code>	Mean and median
<code>min()</code> , <code>max()</code>	Minimum and maximum
<code>std()</code> , <code>var()</code>	Standard deviation and variance
<code>mad()</code>	Mean absolute deviation
<code>prod()</code>	Product of all items
<code>sum()</code>	Sum of all items

GroupBy

```
df = pd.DataFrame({'A': [1, 2, 3],  
                  'B': [3, 4, 5]})
```

	key	data1	data2
0	A	0	2
1	B	1	3
2	A	2	4
3	B	3	5

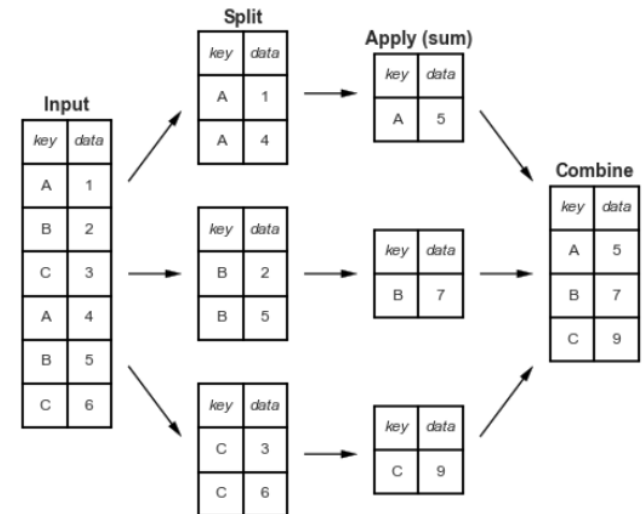
```
df.groupby('key').sum()
```

	key	data1	data2
	A	2	6
	B	4	8

```
df.groupby('key').aggregate(['min', 'max'])
```

		data1		data2	
	key	min	max	min	max
	A	0	2	2	4
	B	1	3	3	5

```
df.groupby('key').apply(your_own_function)
```



Notebook and Exercise time!

Notebook


See 03.08-Aggregation-and-Grouping.ipynb

Exercise time!

See ...



Working with Strings



Working with strings

1. Vectorized String Operations
2. Introducing Pandas String Operations
3. Tables of Pandas String Methods
 - Methods similar to Python string methods
 - Methods using regular expressions
 - Miscellaneous methods
 - Vectorized item access and slicing
 - Indicator variables

GroupBy

String methods

<code>len()</code>	<code>lower()</code>	<code>translate()</code>	<code>islower()</code>
<code>ljust()</code>	<code>upper()</code>	<code>startswith()</code>	<code>isupper()</code>
<code>rjust()</code>	<code>find()</code>	<code>endswith()</code>	<code>isnumeric()</code>
<code>center()</code>	<code>rfind()</code>	<code>isalnum()</code>	<code>isdecimal()</code>
<code>zfill()</code>	<code>index()</code>	<code>isalpha()</code>	<code>split()</code>
<code>strip()</code>	<code>rindex()</code>	<code>isdigit()</code>	<code>rsplit()</code>
<code>rstrip()</code>	<code>capitalize()</code>	<code>isspace()</code>	<code>partition()</code>
<code>lstrip()</code>	<code>swapcase()</code>	<code>istitle()</code>	<code>rpartition()</code>

Notebook and Exercise time!

Notebook

See 03.10-Working-With-Strings.ipynb

Exercise time!

See 3.10_EX_strings.ipynb



Reading Files

Reading files

1. Reading Data

2. Reading CSV files and working with a dataframe

3. Categorical Variables

Reading files

reading csv file

```
data = pd.read_csv('file_name')
```

```
data = pd.read_csv('file_name', sep=';') # separator is ;
```

```
data = pd.read_csv('file_name', sep='; ', decimal=',') # decimal point is ,
```

```
data = pd.read_csv('file_name', names=['n1', 'n2']) # if header is not in file
```

categorical variables

```
bloodtype = pd.Categorical(values, categories=['O-', 'O+', 'B-', 'B+', 'A-', 'A+', 'AB-', 'AB+'])
```

define columns as categorical

```
laptops = pd.read_csv('laptops.csv', dtype={'cpu': 'category', 'brand': 'category'})
```


Notebook and Exercise time!

Notebook

See 03.XTR_ReadingFiles.ipynb

Exercise time!

See ...